




# Prism Games 3.0

**Stochastic Game Verification with Concurrency and Equilibria**

Marco Lavorini  
CMCS 24/25

# Why Probabilistic Verification on Games?

- Agents operate in a noisy, stochastic environment
- Systems increasingly involve concurrently acting agents
- Game-Theory approaches can be used to reason about the competitive or collaborative behaviour of multiple rational agents



*“Program testing can be used to show the presence of bugs, but never to show their absence”*

*-Dijkstra*

# What is a Game?

- A game can be formally defined as a tuple


$$\mathbf{G} = \langle N, S, u, R \rangle$$

- $N = \{1, \dots, n\}$  finite set of players
- $S_i$  set of strategies for player  $i \in N$ ,  $S$  is the set for all players.
- $u : S \rightarrow \mathbb{R}$  payoff function
- $R$  set of possible rules

# More about Games



- The same game can be represented in many way, for example a matrix or a graph.
- There are many types of games:
  - pure, mixed, stochastic, ...
  - Zero-sum, bi-games, ...
- A game can be competitive, collaborative, or even both.

I/II	not confess	confess
not confess	(2,2)	(7,0)
confess	(0,7)	(5,5)

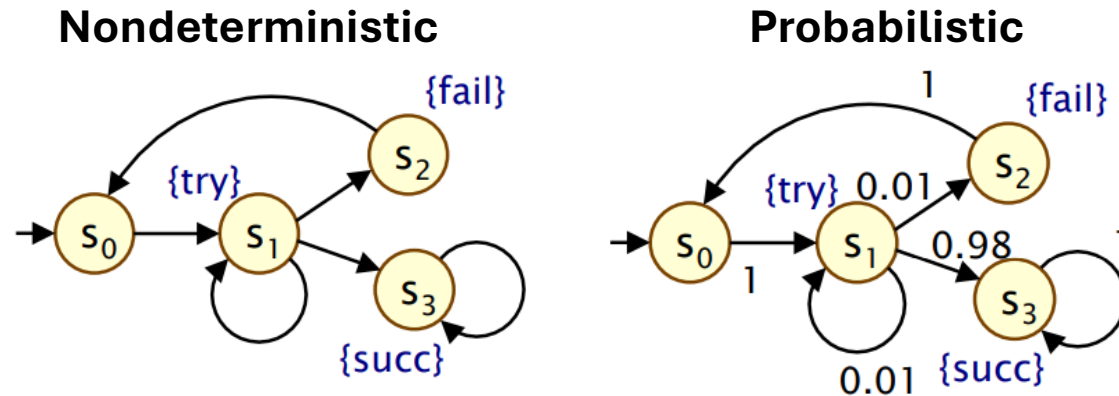


# From Simple Models to Games

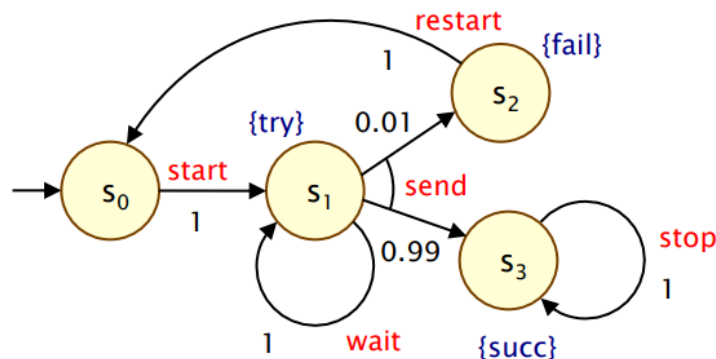
- Labelled Transition System
  - We use CTL
  - Example:
    - ❖  $E [ \text{fail} \cup \text{succ} ]$   
there exists a path that reaches a fail state until it first reaches succ
- Discrete-Time Markov Chain
  - We use PCTL
  - Example:
    - ❖  $P \geq 0.95 [ F \text{succ} ]$   
the probability of reaching succ is at least 95%



# Markov Decision Process



## Probabilities and nondeterminism



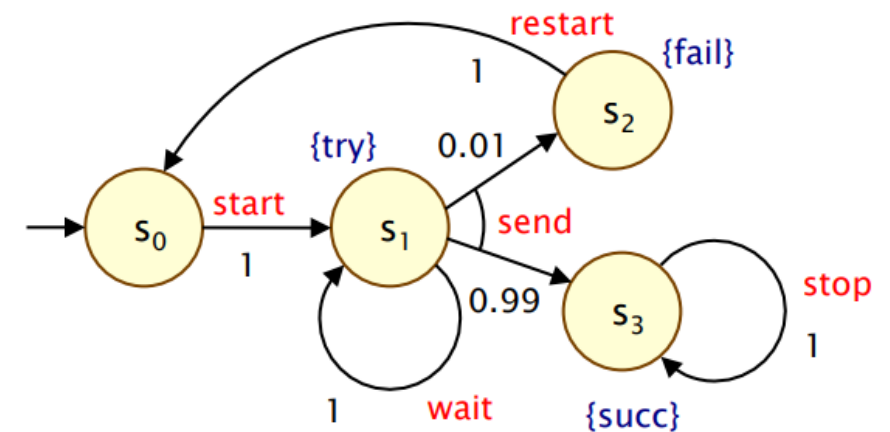
- In each state of the MDP, we have a nondeterministic choice between several discrete probability distributions over the successor states

Example taken from: **Slide 12 of “Lecture 12 – Markov Decision Processes”** from Dave Parker’s *Probabilistic Model Checking* course at Oxford.

# Strategies in an MDP



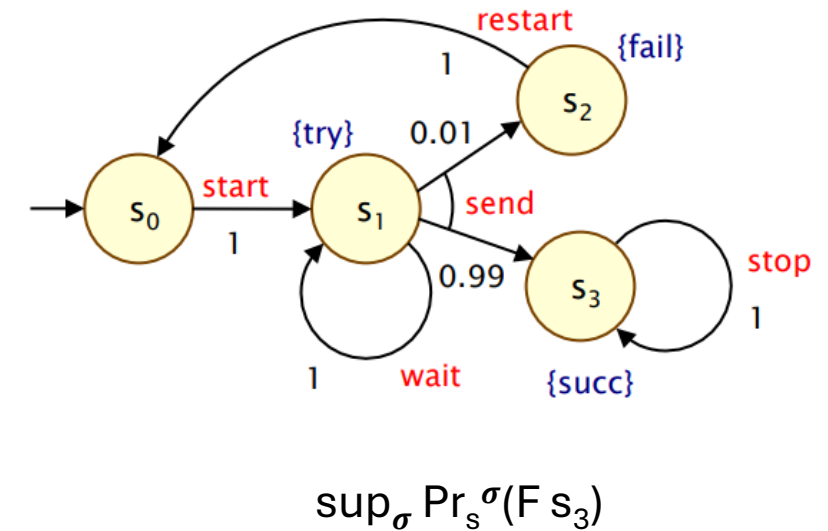
- In any given state we can choose an action
- The choice is still probabilistic
- A series of actions is called a **Strategy**  $\sigma$
- An **optimal strategy** is a strategy that maximise probability
- Example:
  - ❖ What is the maximum probability of reaching state  $s_3$ , achievable by any strategy  $\sigma$ ?
  - ❖  $\sup_{\sigma} \Pr_s^{\sigma}(F s_3)$   
"The supremum over all the possible strategies of the probability from a particular state  $s$  under that strategy  $\sigma$ "



# Value Iteration

- Value Iteration is a dynamic programming approach used for probabilistic model checking


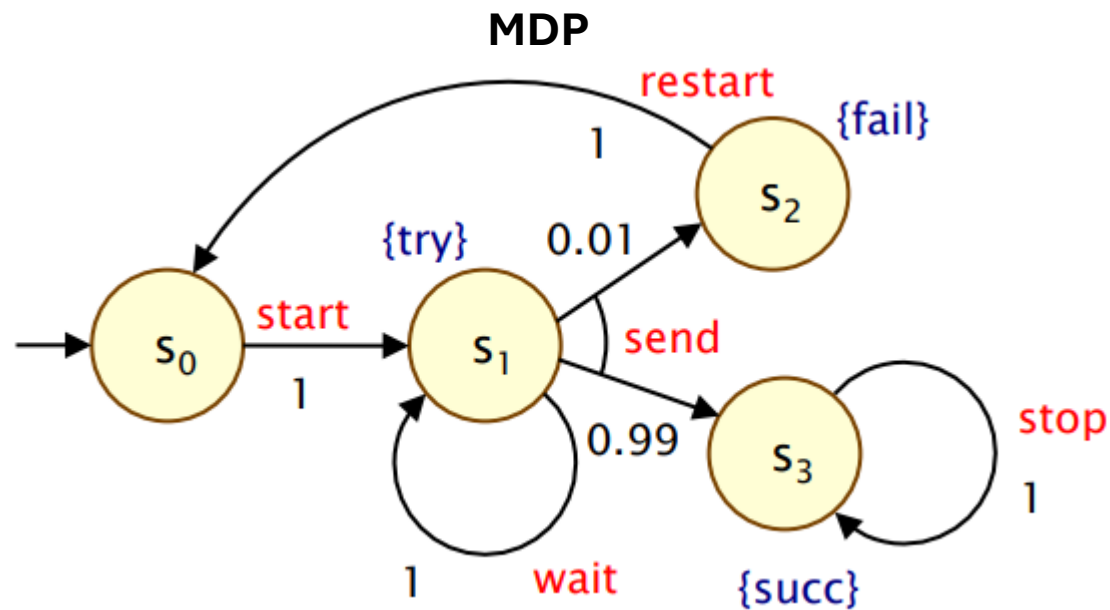
$$p(s) = \begin{cases} 1, & \text{if } s = \text{goal} \\ \max_a \sum_{s'} \delta(s, a)(s') \cdot p(s'), & \text{otherwise} \end{cases}$$





# Stochastic Multiplayer Games

- If we add Players to a MDP, we can easily obtain a Turn-based Stochastic Game. Where each player control a subset of the states set.




# Prism Games

- To model a game on Prism, we will define the **rPATL**,  
«reward probabilistic alternating temporal logic».  
Based on the computational tree logic (**CTL**), it is extended with:
  - Coalition operator  $\langle\langle C \rangle\rangle$  from Alternating-time Temporal Logic (ATL)
  - Probability operator **P** from pCTL
  - Reward operator **R** from PRISM
- Example:
  - ❖  $\langle\langle \text{player}_1, \text{player}_2 \rangle\rangle P_{\geq 90} [F^{\leq 5} \text{goal}_1] \ \& \ R\{\text{"cost"}\} \leq 15 [F \text{goal}_1]$   
“player 1 and 2 have a strategy to ensure that the probability of reaching the  $\text{goal}_1$  state within 5 steps is at least 90%. The expected cost incurred before that happens is **no more than 15.**”

# Coalition for Zero-Sum games

- The coalition operator  $\langle\langle C \rangle\rangle$  from ATL, allows us to reduce every turn-based game into a two player Zero-Sum Game for the desired property, the coalition  $C$  againsts the coalition of every other player.
- What we want to solve now is  


$$\diamond \sup_{\sigma_1} \inf_{\sigma_2} \Pr_s^{\sigma_1 \sigma_2} (F s')$$



# Value Iteration for TSG

- Similar recursive expression used to solve MDP, but now we have to separate the states between the players

$$p(s) = \begin{cases} 1, & \text{if } s = \text{goal} \\ \max_a \sum_{s'} \delta(s, a)(s') \cdot p(s'), & \text{if } s \neq \text{goal and } s \in S_1 \\ \min_a \sum_{s'} \delta(s, a)(s') \cdot p(s'), & \text{if } s \neq \text{goal and } s \in S_2 \end{cases}$$



$$\sup_{\sigma_1} \inf_{\sigma_2} \Pr_s^{\sigma_1 \sigma_2}(F \text{ s}')$$




# Concurrent Stochastic Games

- In a Turn-based Stochastic Game, only **one player at the time** can make an action
- Each player **knows the other player current state** and makes a decision based on that
- A more realistic model have all its agent operating **concurrently**, making strategical choice indipendently


# Concurrent Stochastic Games



Turn-based Stochastic Games (TSG)





Concurrent Stochastic Games (CSG)



# Concurrent Stochastic Games

- CSG are a generalization of TSG, in which we jointly determine the probabilistic successor state.
- In TSG the best strategy was deterministic, now the best strategy can be randomized
- The logic used to model CSG properties is still rPATL but with some changes





# Value Iteration for CSG

- What we want to solve is still:  
    ❖  $\sup_{\sigma_1} \inf_{\sigma_2} \Pr_s^{\sigma_1 \sigma_2}(F \mid s')$
- And the recursive expression is now:

$$p(s) = \begin{cases} 1, & \text{if } s = \text{goal} \\ \text{val}(G), & \text{otherwise} \end{cases}$$

Where  $G$  is the bimatrix game with

$$g_{ij} = \sum_{s'} \delta(s, (a_i, b_j)(s')) \cdot p(s')$$

I/II	not confess	confess
not confess	(2, 2)	(7, 0)
confess	(0, 7)	(5, 5)

- The bimatrix game can be solved by Linear Programming



# Modelling CSG



- Players are now associated with **modules, instead of states**, this way a state is not bound to a single player.
- Modules represent individual agents' behaviour and choices, modules with no nondeterministic choice (like channels) **do not need to be tied to a player**.
- Player concurrently select between multiple **possible** actions
- This results in a complex, interdependent probabilistic distribution for the resulting joint-actions.

# Medium Access Control in CSG

```

1  csg
2  // Player specification
3  player p1 mac1 endplayer
4  player p2 mac2 endplayer
5  // Max energy per user
6  const int emax;
7  // User 1
8  module mac1
9      s1 : [0..1] init 0; // Has user 1 sent?
10     e1 : [0..emax] init emax; // Energy level of user 1
11     [w1] true -> (s1'=0); // Wait
12     [t1] e1>0 -> (s1'=c'?0:1) & (e1'=e1-1); // Transmit
13 endmodule
14 // Define second user using module renaming
15 module mac2 = mac1 [ s1=s2, e1=e2, w1=w2, t1=t2 ] endmodule

```

Player are associated to modules.

```

1  // Probability qi for transmission success when i users send
2  const double q1;
3  const double q2;
4  // Channel (computes joint transmission probabilities)
5  module channel
6      c : bool init false; // Did a collision occur during transmission?
7      [t1,w2] true -> q1:(c'=false) + (1-q1):(c'=true); // User 1 transmits
8      [w1,t2] true -> q1:(c'=false) + (1-q1):(c'=true); // User 2 transmits
9      [t1,t2] true -> q2:(c'=false) + (1-q2):(c'=true); // Both transmit
10 endmodule

```

**Transition** and **Guarded Command** are now a list since the players move concurrently.

Modules with no nondeterministic choice do not need to be tied to a player.

```

1  // Reward structures
2  rewards "mess1" // Number of messages sent by user 1
3      s1=1 : 1;
4  endrewards
5  rewards "mess2" // Number of messages sent by user 2
6      s2=1 : 1;
7  endrewards
8  rewards "send2" // Number of times users 1 and 2 transmit simultaneously
9      [t1,t2] true : 1;
10 endrewards

```



# From Zero-Sum to Equilibria

- With rPATL we solved **zero-sum** game with the maximisation of the probability of reaching the wanted goal for a coalition.
- In Prism game 3.0 we can compute the ***social-welfare optimal Nash equilibria*** in which we maximise (or minimize) the sum of the values associated to the objectives for each player.
- The general idea is that player can have distinct objective which are not directly opposing (zero-sum)
- To do this we add to rPATL the **+** operator



# What is an equilibria?

- Let  $G = (N, S, \mu)$  be a Game.
- A Nash Equilibrium is a strategy profile  $x^* \in S$  such that the strategy  $x_i^*$  is a best response to the strategy profile  $x_{-i}^*$  for all  $i \in N$

$$u_i(x_i^*, x_{-i}^*) \geq u_i(x_i, x_{-i}^*) \quad \text{for all } x_i \in S_i$$

I/II	not confess	confess
not confess	(2,2)	(7,0)
confess	(0,7)	(5,5)

(confess, confess) is the unique Nash Equilibrium

Not Pareto optimal: both player could get less years

Not socially optimal: it is not the best joint best result

# Equilibria II



- The *social-welfare optimal Nash equilibria* is a kind of equilibria where there is no incentive for any player to unilaterally change strategy, given the maximisation of the **sum** of the players' payoff.

**Zero-sum**

$$\langle\langle\{\text{player}_1\}\rangle\rangle_{\max=?} \mathbf{P}[F^{\leq k} \text{goal}_1]$$



**SWNE**

$$\langle\langle\{\text{player}_1, \text{player}_2\}\rangle\rangle_{\max=?} (\mathbf{P}[F^{\leq k} \text{goal}_1] + \mathbf{P}[F^{\leq k} \text{goal}_2])$$

# Value Iteration for CSG with equilibria

- The recursive expression is now:

$$p(s) = \begin{cases} (1,1), & \text{if } s_1 = goal_1, s_2 = goal_2 \\ (P_{max}(s, goal_1), 1), & \text{if } s_1 \neq goal_1, s_2 = goal_2 \\ (1, P_{max}(s, goal_2)), & \text{if } s_1 = goal_1, s_2 \neq goal_2 \\ SWNE(G(s)), & \text{if } s_1 \neq goal_1, s_2 \neq goal_2 \end{cases}$$

- Now we have pairs of value in order to consider both player objectives
- The first three cases are a MDP

# Equilibria for unbounded properties


- ‘Unbounded properties’ are properties that do not specify a time or step limit. They require reasoning about behavior potentially extending indefinitely into the future
  - ❖ Example:  
“probability of eventually reaching state **s**’ (regardless of the number of steps).
- Nash Equilibria might not exist or might not be computable, hence PRISM-games computes  **$\epsilon$ -Nash equilibria**
  - A strategy profile  $\sigma = (\sigma_1, \dots, \sigma_n)$  for a CSG is a  **$\epsilon$ -NE** for state **s** and objective  $X_1, \dots, X_n$  iff:

$$\Pr_s^\sigma(X_i) \geq \sup_{\sigma'_i} \Pr_s^{(\sigma_{-i}, \sigma'_i)}(X_i) - \epsilon, \quad \forall i.$$


# Case studies



## Future markets investor



## Robot coordination






# Future Market Investors





- A stock market (the model) evolves stochastically
- Two investors decide when to invest
- The market can refuse the investments
- Making concurrent decision is more realistic, the market can not observe the investors strategies.



(a) **Future markets investor:** avoiding unrealistic strategy choices using CSGs

# Robot coordination


- Two robots navigating in a  $l \times l$  grid
- Starting from the opposite corners, the goal is to reach the opposite corners.
- Each step has a probability  $q$  to fail to simulate an obstacle
- When robots collide, they will need to wait before moving



(b) **Robot coordination:** using equilibria for mutually beneficial navigation plans

# Example: Rock Paper Scissor

- Very intuitive game with simple rules.
- Zero-sum game with no *Pure strategy Nash Equilibria*
- By modelling the game in Prism 3.0 we can observe the expected probability of the simple model, and calculate the *Mixed strategies Nash Equilibria*





# DEMO



# Example: Rock Paper Scissor

- “What is the maximum probability that Player 1 wins before ever losing?”

**<<p1>>Pmax=? [ win!=2 U win=1 ]**

$$p = \frac{1}{3} \cdot 1 \text{ (win)} + \frac{1}{3} \cdot 0 \text{ (loss)} + \frac{1}{3} \cdot p \text{ (draw, retry)}$$

$$p = \frac{1}{3} + \frac{1}{3}p \implies p - \frac{1}{3}p = \frac{1}{3} \implies \frac{2}{3}p = \frac{1}{3} \implies p = \frac{1}{2}$$

- 
- “What is the maximum probability that Player 1 gets exactly 1 win and exactly 2 draws in 3 rounds?”

**<<p1>>Pmax=? [ F (win\_count=1 & draw\_count=2) ]**

$$3 \times \frac{1}{27} = \frac{3}{27} = \frac{1}{9}$$